


TV on $F \cup \{\perp\}$ from the GPU double-precision output

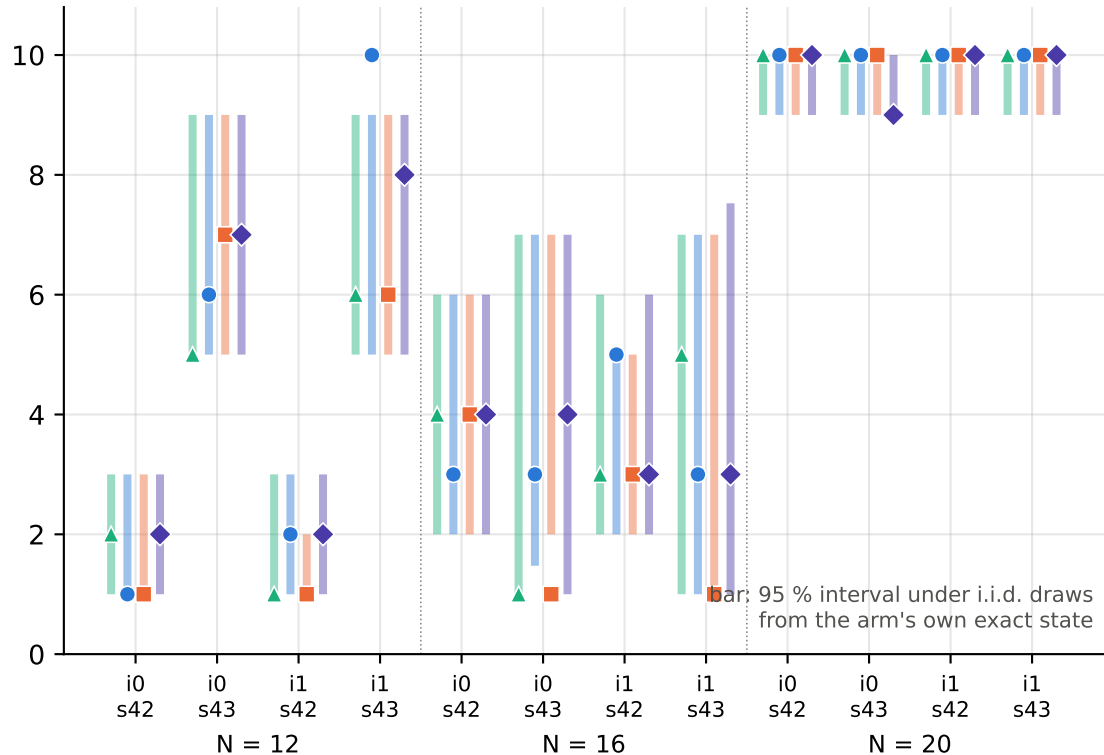
Same angles, different backend: the output distribution



- CPU, double (1 thread)
- GPU, single
- ◆ SV1, double (cloud)
- single-precision ϵ
- double-precision ϵ

Distinct best-feasible portfolios in 10 batches of 1,000 shots

Same output distribution, ten batches: the answer



- ▲ GPU, double (reference)
- CPU, double (1 thread)
- GPU, single
- ◆ SV1, double (cloud)

Bar: 95 % interval under i.i.d. draws from the arm's own exact state